

Zhuofeng Wu

3442 North Quad, 105 S. State St., Ann Arbor, MI, USA, 48109-1285
(E) zhuofeng@umich.edu (C) (734) 780-9664 (W) <https://cserxy.github.io/>

EDUCATION

University of Michigan, Ann Arbor, US

Aug 2018 – present

Ph.D. candidate in School of Information (Advisor: V.G. Vinod Vydiswaran)
Natural Language Processing, Machine Learning

Zhejiang University, Hangzhou, China

Sept 2013 - Jun 2017

B.E. in Computer Science (Overall GPA: 3.82/4.0, Top 5% among all 215 students)
Pursuit Science Class, Chu Kochen Honors College (CKC College)
Received waiver for the National College Entrance Exam to enter Zhejiang University from **1st Prize in National Olympiad in Informatics in Provinces** (top 1.8% over 60,000 participants)

PUBLICATIONS

HiCL: Hierarchical Contrastive Learning of Unsupervised Sentence Embeddings

Zhuofeng Wu, Chaowei Xiao, V. G. Vydiswaran
In Proceedings of Findings EMNLP 2023.

PLANNER: Generating Diversified Paragraph via Latent Language Diffusion Model

Yizhe Zhang, Jiatao Gu, Zhuofeng Wu, Shuangfei Zhai, Josh Susskind, Navdeep Jaitly
In Proceedings of NeurIPS 2023. ([pdf](#))

Defending against Insertion-based Textual Backdoor Attacks via Attribution

Jiazha Li, Zhuofeng Wu, Wei Ping, Chaowei Xiao, V. G. Vydiswaran
In Proceedings of Findings ACL 2023. ([pdf](#))

IDPG: An Instance-Dependent Prompt Generation Method

Zhuofeng Wu, Sinong Wang, Jiatao Gu, Rui Hou, Yuxiao Dong, V. G. Vydiswaran, Hao Ma
In Proceedings of NAACL 2022 (**Oral Presentation**). ([pdf](#)) ([video](#))

Identify Shifts of Word Semantics through Bayesian Surprise

Zhuofeng Wu, Cheng Li, Zhe Zhao, Fei Wu, Qiaozhu Mei
In Proceedings of SIGIR 2018 (**Oral Presentation**). ([pdf](#))

PREPRINT

Chatgpt as an attack tool: Stealthy textual backdoor attack via blackbox generative model trigger

Jiazha Li, Yijin Yang, Zhuofeng Wu, V. G. Vydiswaran, Chaowei Xiao
arXiv preprint arXiv:2304.14475 (In submission to EACL 2024) ([pdf](#))

Clear: Contrastive learning for sentence representation

Zhuofeng Wu, Sinong Wang, Jiatao Gu, Madian Khabsa, Fei Sun, Hao Ma
arXiv preprint arXiv:2012.15466 (2020). ([pdf](#))

EXPERIENCE

Apple Machine Learning Research

Apr 2023 – Aug 2023

Research Intern, Advisor: Dr. Yizhe Zhang

Knowledge Distillation from LLM to Small Models: A Perspective from Question Decomposition

- Leverage LLMs such as GPT-4 to decompose a question into several related sub-questions.
- Fine-tune a model on the generated question-subquestions pair initially, and further train it based on the rewards from GPT-4.
- Extensive evaluations on GSM8k and DROP dataset show our proposed method can catch LLMs' question decomposition capability (and sometimes even better, e.g., better than ChatGPT).
- This work is an ongoing project and targets at **ICML'24**.

Facebook AI

Research Intern, Advisor: Dr. Sinong Wang

May 2021 – Aug 2021

IDPG: An Instance-Dependent Prompt Generation Method

- First customized prompt for each input rather than one prompt for all inputs.
- Offered comparable performance to Adapter-based methods while using fewer parameters.

- Extensive evaluations on ten natural language understanding tasks show that IDPG consistently outperforms task-specific prompt tuning methods by 1.6–3.1 points.
- This work was presented at **NAACL’22** as **oral**.

Research Intern, Advisor: Dr. Sinong Wang

May 2020 – Aug 2020

CLEAR: Contrastive Learning for Sentence Representation

- Proposed to align the representation of different argumentation for same sentence.
- Explored several argumentations and their combinations in the text domain.
- Revealed that different argumentations in pre-training enhance the model’s different abilities.
- Outperformed several baselines (including BERT & RoBERTa) on GLUE & SentEval benchmark.

Alibaba Group

May 2019 – Aug 2019

Research Intern, Advisor: Dr. Fei Sun

Seg-BERT: A Hierarchical Structure for Document Classification

- Applied a hierarchical structure for the long text classification.
- Outperformed the state-of-the-art by a large margin on IMDB.
- Proposed to mask sentence in pre-training to improve the performance.

School of Information, University of Michigan

Aug 2018 – Nov 2020

Research Assistant, Advisors: Prof. Qiaozhu Mei, Prof. Daniel Romero

Relocation Detection with Extra Information from Online Social Behavior on Twitter

- Proposed to extract extra information from online social behavior to help the relocation detection.

School of Information, University of Michigan

Apr 2016 – Apr 2018

Research Intern, Advisor: Prof. Qiaozhu Mei

Identify Shifts of Word Semantics through Bayesian Surprise

- Explicitly established the stable topological structure of word semantics and identified the surprising changes over time.
- Proposed a statistical framework to apply **Bayesian Surprise** in detecting the meaning-changed words in **temporal-based word semantic networks**. This framework can be generalized to finding the change points in many other networks.
- Conducted experiments on ACMDL, DBLP and Google Books Ngram data set for synthetic evaluation which artificially introducing changes to a corpus. Outperformed the state-of-the-art by a large margin.
- This work was presented at **SIGIR’18** as **oral** and was adopted as a part of a **KDD’18 Workshop Keynote Talk** “Identifying Shifts in Evolutionary Semantic Spaces”.

A Tool to Visualize the Evolution of Conference Topics

- Visualized a 40-year evolution of data science related communities and embedded papers, keywords, authors in the same space.
- Provided a powerful tool for researchers to model the research focus of different conferences.
- This work was presented in an invited talk in **KDD’18 Deep Learning Day** by Prof. Mei.

Digital Media Computing & Design Lab, Zhejiang University

Sept 2014 - Mar 2016

Research Assistant, Advisor: Prof. Fei Wu

Explored how to train different embedding models and implemented several word representation algorithms in C++.

SKILLS

Programming Languages: C, C++, Python, Verilog, Pascal

Frameworks & Tools: PyTorch, Fairseq, LaTeX, Vim, Git

SERVICE

Conference Reviewer: ACL’23, SIGIR’23, EMNLP’23, NeurIPS’23.

ACL Rolling Reviewer: Dec’22, Apr’23, Jun’23.

Student volunteer for SIGIR’18, NAACL’22.

GSI for SI 670 Applied Machine Learning, SI 630 Natural Language Processing, SI 650 Information Retrieval, LHS 712

Natural Language Processing for Health.

AWARDS

SIGIR Student Travel Grant, 2018.

Outstanding Graduates of Zhejiang Province, 2017.

2nd Prize of Excellent Undergraduate Scholarship, 2014.

3rd Prize in Collegiate Programming Contest of Zhejiang University, 2014, 2015.

1st Prize in National Olympiad in Informatics in Provinces in 2012.

1st Prize in National Olympiad in Mathematics in Provinces in 2010.